


SIU AI API — End-User Guide

Gateway: [SIU AI API gateway](#)

Hosted model: SIU's Qwen 3.8 27B deployment

API style: OpenAI compatible through LiteLLM

Guide version: 2026-08-21



SIU AI API connection settings overview.

The SIU Computer Science AI API can be used from ordinary programs, REST clients, IDE assistants, and coding agents. Most integrations require only three values:

Setting	Value
OpenAI-compatible base URL	<code>`<a href="https://api.cs.siu.edu/v1">https://api.cs.siu.edu/v1</a>`</code>
API key	Your individually assigned SIU AI key
Model ID	The exact Qwen model ID returned by <code>`GET /v1/models`</code>

Do not guess the model ID. The model's display name is Qwen 3.8 27B, but the public LiteLLM alias may be different. Run the model-discovery request first and copy the exact `id` value.

Protect your key. Treat it like a password. Do not put it in source code, screenshots, assignments, public repositories, shared configuration files, or browser-side JavaScript.

Contents

- [Choose the right client](#)
- [Get ready and discover the model](#)
- [cURL](#)
- [Python OpenAI SDK](#)
- [JavaScript / TypeScript OpenAI SDK](#)
- [Continue for VS Code or JetBrains](#)
- [Cline for VS Code](#)
- [OpenCode](#)
- [Aider](#)
- [Postman](#)
- [LangChain and LangGraph](#)
- [Claude Code through LiteLLM](#)
- [Codex CLI](#)
- [Kilo Code](#)
- [Roo Code](#)
- [Troubleshooting](#)
- [Security and responsible use](#)
- [Included files](#)

1. Choose the right client

- **cURL:** Best for first tests and troubleshooting. Routes: `/v1/models` and `/v1/chat/completions`. Start here.

- **Python OpenAI SDK:** Best for Python assignments and services. Route: ``/v1/chat/completions``. Recommended.
- **JavaScript OpenAI SDK:** Best for Node.js and TypeScript applications. Route: ``/v1/chat/completions``. Recommended.
- **Continue:** Best for IDE chat, code edits, and optional agent mode. Route: ``/v1/chat/completions``. Recommended; validate tools before agent mode.
- **Cline:** Best for agentic development in VS Code. Route: ``/v1/chat/completions``. Recommended with command approval enabled.
- **OpenCode:** Best for a terminal coding agent. Route: ``/v1/chat/completions``. Recommended.
- **Aider:** Best for Git-oriented terminal coding. Route: ``/v1/chat/completions``. Recommended.
- **Postman:** Best for visual REST testing and instruction. Routes: ``/v1/models`` and ``/v1/chat/completions``. Recommended.
- **LangChain / LangGraph:** Best for AI application development. Route: ``/v1/chat/completions``. Recommended.
- **Kilo Code:** Best for a VS Code coding assistant. Routes: ``/v1/models`` and ``/v1/chat/completions``. Recommended; validate tools.
- **Roo Code:** Best for a VS Code coding agent with a LiteLLM provider. Routes: ``/v1/model/info`` and chat routes. Recommended; use the fallback if discovery is disabled.
- **Claude Code:** Provides the Claude Code interface routed to Qwen. Route: ``/v1/messages``. Advanced and compatibility-dependent.
- **Codex CLI:** Provides the Codex interface routed to Qwen. Route: ``/v1/responses``. Advanced; publish only after endpoint validation.

A practical starting set for students: cURL, Python, JavaScript, Continue, Cline, OpenCode, Aider, and Postman. Claude Code and Codex should remain in an advanced section because they use different wire protocols and may not reproduce all native vendor behavior when backed by Qwen.

2. Get ready and discover the model

REQUIRED FIRST STEP

Find the public model ID first

Do not guess the model alias—ask the gateway what your key can access.

1

Set your key

Use a temporary shell variable; do not paste it into screenshots.

2

List models

Call the OpenAI-compatible models endpoint.

3

Copy the ID

Use the exact Qwen 3.8 27B ID in every client.

```
Terminal — Linux, macOS, or WSL

export SIU_AI_KEY="paste_your_assigned_key_here"

curl https://api.cs.siu.edu/v1/models \
-H "Authorization: Bearer $SIU_AI_KEY"

# Optional: pretty-print with jq
curl -s https://api.cs.siu.edu/v1/models \
-H "Authorization: Bearer $SIU_AI_KEY" | jq .
```

Read the response

Look inside the data array.

```
response.json

{
  "data": [
    { "id": "MODEL_ID" }
  ]
}
```

✓

Save this value

Set SIU_AI_MODEL to the exact ID.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Finding and copying the public model ID.

2.1 Obtain an API key

Use the key issued to you by the department. Keys may have individual model access, rate limits, concurrency limits, and budgets. Do not use another person's key.

2.2 Set temporary environment variables

Linux, macOS, or WSL - set variables

```
export SIU_AI_BASE_URL="https://api.cs.siu.edu/v1"
export SIU_AI_KEY="paste_your_assigned_key_here"
```

Windows PowerShell - set variables

```
$env:SIU_AI_BASE_URL = "https://api.cs.siu.edu/v1"
$env:SIU_AI_KEY = "paste_your_assigned_key_here"
```

These values last for the current terminal session. That is usually safer than permanently storing a key in shell startup files.

2.3 List the models your key can use

Linux, macOS, or WSL - list models

```
curl --fail-with-body --silent --show-error \  
"$SIU_AI_BASE_URL/models" \  
-H "Authorization: Bearer $SIU_AI_KEY"
```

With `jq` installed:

```
curl -s "$SIU_AI_BASE_URL/models" \  
-H "Authorization: Bearer $SIU_AI_KEY" \  
| jq -r '.data[].id'
```

Windows PowerShell - list models

```
$headers = @{ Authorization = "Bearer $env:SIU_AI_KEY" }  
$response = Invoke-RestMethod \  
-Uri "$env:SIU_AI_BASE_URL/models" \  
-Headers $headers \  
-Method Get  
  
$response.data.id
```

Copy the exact ID corresponding to the SIU Qwen 3.8 27B deployment, then set it:

```
export SIU_AI_MODEL="the_exact_id_returned_by_v1_models"  
$env:SIU_AI_MODEL = "the_exact_id_returned_by_v1_models"
```

A ready-made discovery script is included at `examples/curl/list-models.sh`.

3. cURL

API QUICK START

Test with cURL

The fastest way to verify the URL, key, model, and request body.

1 Set variables
Keep the secret out of shell history when possible.

2 Send POST
Use /v1/chat/completions and JSON.

3 Read content
The answer is in choices[0].message.content.

Terminal — chat completion

```
export SIU_AI_MODEL="MODEL_ID_FROM_V1_MODELS"  
  
curl https://api.cs.siu.edu/v1/chat/completions \  
-H "Authorization: Bearer $SIU_AI_KEY" \  
-H "Content-Type: application/json" \  
-d '{  
  "model": "MODEL_ID_FROM_V1_MODELS",  
  "messages": [  
    {"role": "user", "content": "Explain recursion."}  
  ],  
  "temperature": 0.2  
'
```

Expected result
HTTP 200 and a choices array.

200OK

response excerpt

```
"choices": [{  
  "message": {  
    "content": "..."  
  }  
}]
```


Included file: examples/curl/chat.sh

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Complete cURL chat-completion request and response.

Use cURL before configuring another application. A successful request confirms that the network path, TLS connection, API key, model access, and basic LiteLLM route all work.

3.1 Send a chat-completion request

The included script safely constructs the JSON from environment variables:

```
cd examples/curl
./chat.sh
```

The equivalent request is:

```
curl --fail-with-body --silent --show-error \
  "$SIU_AI_BASE_URL/chat/completions" \
  -H "Authorization: Bearer $SIU_AI_KEY" \
  -H "Content-Type: application/json" \
  --data "$(cat <<JSON
{
  "model": "$SIU_AI_MODEL",
  "messages": [
    {
      "role": "system",
      "content": "You are a concise and careful programming assistant."
    },
    {
      "role": "user",
      "content": "Write a Python function that checks whether a string is a
palindrome."
    }
  ],
  "temperature": 0.2,
  "stream": false
}
JSON
)"
```

A successful response normally contains:

```
{
  "choices": [
    {
      "message": {
        "role": "assistant",
        "content": "..."
      }
    }
  ]
}
```

3.2 PowerShell request

```

$headers = @{
    Authorization = "Bearer $env:SIU_AI_KEY"
    "Content-Type" = "application/json"
}

$body = @{
    model = $env:SIU_AI_MODEL
    messages = @(
        @{
            role = "user"
            content = "Explain recursion using a small Python example."
        }
    )
    temperature = 0.2
    stream = $false
} | ConvertTo-Json -Depth 5

$response = Invoke-RestMethod `
    -Uri "$env:SIU_AI_BASE_URL/chat/completions" `
    -Headers $headers `
    -Method Post `
    -Body $body

$response.choices[0].message.content

```

4. Python OpenAI SDK

PYTHON

Use the Python OpenAI SDK

Change the client base URL; the rest looks like a normal OpenAI-compatible request.

1 Create an environment

python -m venv .venv, then activate it.

2 Install the SDK

python -m pip install --upgrade openai

3 Export settings

Set SIU_AI_KEY and SIU_AI_MODEL.

quickstart.py

```

import os
from openai import OpenAI

client = OpenAI(
    base_url="https://api.cs.siu.edu/v1",
    api_key=os.environ["SIU_AI_KEY"],
)

reply = client.chat.completions.create(
    model=os.environ["SIU_AI_MODEL"],
    messages=[
        {"role": "user", "content": "Explain stacks."}
    ],
    temperature=0.2,
)
print(reply.choices[0].message.content)

```

Run it

Cross-platform commands

Terminal

```

python -m pip install openai
export SIU_AI_KEY="..."
export SIU_AI_MODEL="..."
python quickstart.py

```

Streaming included

Use examples/python/streaming.py for token-by-token output.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Python OpenAI SDK configuration.

LiteLLM accepts OpenAI-compatible requests, so Python applications can use the standard `OpenAI` client with a custom `base\_url`.

4.1 Create an environment and install the SDK

Linux, macOS, or WSL - Python environment

```
cd examples/python
python3 -m venv .venv
source .venv/bin/activate
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

Windows PowerShell - Python environment

```
cd examples\python
py -m venv .venv
.\.venv\Scripts\Activate.ps1
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

4.2 Run the included example

Make sure `SIU\_AI\_KEY` and `SIU\_AI\_MODEL` are set, then run:

```
python quickstart.py
```

The essential code is:

```
import os
from openai import OpenAI

client = OpenAI(
    base_url="https://api.cs.siu.edu/v1",
    api_key=os.environ["SIU_AI_KEY"],
)

response = client.chat.completions.create(
    model=os.environ["SIU_AI_MODEL"],
    messages=[
        {
            "role": "user",
            "content": "Explain Python list comprehensions with one short
example.",
        }
    ],
    temperature=0.2,
)

print(response.choices[0].message.content)
```


4.3 Stream output as it is generated

```
python streaming.py
```

The complete examples include timeouts, retries, and basic error handling:

- `examples/python/quickstart.py`
- `examples/python/streaming.py`

5. JavaScript / TypeScript OpenAI SDK

NODE.JS

Use the JavaScript / TypeScript SDK

Point the official OpenAI client at the SIU gateway with baseURL.

- 1 Initialize**
npm init -y
- 2 Install**
npm install openai
- 3 Run**
node quickstart.mjs

```
import OpenAI from "openai";

const client = new OpenAI({
  baseURL: "https://api.cs.siu.edu/v1",
  apiKey: process.env.SIU_AI_KEY,
});

const reply = await client.chat.completions.create({
  model: process.env.SIU_AI_MODEL,
  messages: [
    { role: "user", content: "Explain promises." },
  ],
  temperature: 0.2,
});

console.log(reply.choices[0].message.content);
```

Environment

Keep secrets outside source code.

```
export SIU_AI_KEY="..."
export SIU_AI_MODEL="..."
npm install
npm start
```

Do not bundle keys
Never place an API key in browser-side JavaScript.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

JavaScript OpenAI SDK configuration.

Use this approach for Node.js, server-side TypeScript, Electron, and other trusted runtimes.

Never place an API key in browser-delivered JavaScript. A key embedded in React, Vue, Angular, or ordinary browser code can be extracted by anyone who loads the application. Put the API call in a trusted server-side component.

5.1 Install and run

This example requires **Node.js 22 or newer**. Check the active version before installing dependencies:

```
cd examples/javascript
node --version
```

```
npm install
npm start
```

The version should begin with `v22` or a higher number. If it is older, upgrade Node.js using your system administrator's approved method, open a new terminal, and run the commands again. A Python virtual environment such as `(.venv)` does not install or change Node.js.

For streaming output:

```
npm run stream
```

The essential code is:

```
import OpenAI from "openai";

const client = new OpenAI({
  baseURL: "https://api.cs.siu.edu/v1",
  apiKey: process.env.SIU_AI_KEY,
});

const response = await client.chat.completions.create({
  model: process.env.SIU_AI_MODEL,
  messages: [
    {
      role: "user",
      content: "Explain async/await in JavaScript with one small example.",
    },
  ],
  temperature: 0.2,
});

console.log(response.choices[0].message.content);
```

Included files:

- `examples/javascript/quickstart.mjs`
- `examples/javascript/streaming.mjs`
- `examples/javascript/package.json`

6. Continue for VS Code or JetBrains

CONTINUE

Configure Continue in VS Code

Use a local YAML configuration and a separate secret file.

CONTINUE

Local Config

config.yaml

.env

```
name: SIU AI
version: 1.0.0
schema: v1

models:
  - name: SIU Qwen 3.8 27B
    provider: openai
    model: MODEL_ID_FROM_V1_MODELS
    apiBase: https://api.cs.siu.edu/v1
    apiKey: "${ secrets.SIU_AI_KEY }"
    roles: [chat, edit, apply]
    useResponsesApi: false

# Add tool_use only after validation.
```

Setup checklist

Continue reads the local config automatically.

- 1 Install Continue
VS Code Extensions → Continue
- 2 Open local config
Continue menu → Configs → gear
- 3 Store the key
~/continue/.env
- 4 Reload
Open chat and select SIU Qwen

Agent mode
Keep command approval on; leave auto-approv...

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Continue configuration and setup checklist.

Continue is a good default IDE guide because it supports chat, code editing, and optional agent capabilities while using an OpenAI-compatible provider.

6.1 Install Continue

Install **Continue** from the extension or plugin marketplace for your IDE. Open Continue, select the configuration menu, and open the local configuration.

Common local paths:

- Linux and macOS: `~/continue/config.yaml`
- Windows: `%USERPROFILE%\continue\config.yaml`
- Secret file: `~/continue/.env` or `%USERPROFILE%\continue\.env`

6.2 Copy the supplied configuration

Copy `examples/continue/config.yaml` to the local Continue configuration path.
Replace:

```
model: replace_with_model_id_from_v1_models
```

with the exact ID returned by `GET /v1/models`.

The relevant configuration is:

```
name: SIU AI
version: 1.0.0
schema: v1

models:
  - name: SIU Qwen 3.8 27B
    provider: openai
    model: replace_with_model_id_from_v1_models
    apiBase: https://api.cs.siu.edu/v1
    apiKey: "${ secrets.SIU_AI_KEY }"
    roles:
      - chat
      - edit
      - apply
    useResponsesApi: false
```

6.3 Store the key separately

Copy `examples/continue/.env.example` to `~/continue/.env`, rename it to `.env`, and replace the placeholder:

```
SIU_AI_KEY=your_assigned_key
```

Reload the IDE, open Continue, and select **SIU Qwen 3.8 27B**.

6.4 Agent mode

Start with ordinary chat and edit requests. Only add the following capability after the department confirms native tool calling works with the deployed model and client version:

```
capabilities:
  - tool_use
```

Keep command approval enabled. Do not enable unrestricted automatic approval on coursework, research code, credentials, or important repositories.

7. Cline for VS Code

CLINE

Configure Cline

Select OpenAI Compatible, enter four values, and verify the connection.

API Provider Settings

Cline → Settings → API Provider

API Provider

OpenAI Compatible

Base URL

https://api.cs.siu.edu/v1

API Key

.....

Model ID

MODEL_ID_FROM_V1_MODELS

Context Window

Leave default unless SIU publishes a value

Max Output

Leave default unless SIU publishes a value

Verify Connection

Safe first run

Coding agents can edit files and run commands.

- 1 Open a disposable test project.
- 2 Keep Auto-approve OFF.
- 3 Ask Cline to explain one file first.
- 4 Review every proposed command.
- 5 Commit or back up before edits.

Tool calling

If tools loop, switch to chat-only mode.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Cline provider settings and first-run checklist.

7.1 Enter the provider settings

Open Cline, select the settings gear, and configure:

Field	Value
API Provider	OpenAI Compatible
Base URL	` <a href="https://api.cs.siu.edu/v1">https://api.cs.siu.edu/v1</a> `
API Key	Your assigned SIU AI key
Model ID	Exact ID returned by `GET /v1/models`
Context Window	`32768`
Max Output Tokens	`4096`

Select **Verify** or **Verify Connection**.

7.2 Perform a safe first test

- Open a small disposable or already-committed project.

- Keep **Auto-approve** off.
- Ask Cline to explain one file without changing it.
- Request one small edit.
- Review the proposed diff and every command before approval.

If chat works but tool calls repeat, appear as plain text, or fail to execute, switch to chat-only behavior and report the issue with the client version and a redacted transcript.

8. OpenCode

OPENCODE

Configure OpenCode

Define a custom OpenAI-compatible provider, then store the key with `/connect`.

opencode.json

```

{
  "$schema": "https://opencode.ai/config.json",
  "model": "siu/MODEL_ID",
  "provider": {
    "siu": {
      "npm": "@ai-sdk/openai-compatible",
      "name": "SIU AI",
      "options": {
        "baseUrl": "https://api.cs.siu.edu/v1"
      }
    },
    "models": {
      "MODEL_ID": { "name": "SIU Qwen 3.8 27B" }
    }
  }
}

```

Connect and select

The provider ID must be exactly siu.

OpenCode

✓ Why /connect?

It keeps the API key out of the project configuration file.

i Included

examples/opencode/opencode.json

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

OpenCode provider configuration and connection workflow.

OpenCode can use a custom provider backed by `@ai-sdk/openai-compatible`.

8.1 Create the provider configuration

Copy `examples/opencode/opencode.json` into the project directory as `opencode.json`. A global configuration can instead be placed in OpenCode's documented user configuration location.

Replace both occurrences of `replace_with_model_id_from_v1_models` with the exact model ID.

```

{
  "$schema": "https://opencode.ai/config.json",

```

```

"model": "siu/replace_with_model_id_from_v1_models",
"provider": {
  "siu": {
    "npm": "@ai-sdk/openai-compatible",
    "name": "SIU AI",
    "options": {
      "baseUrl": "https://api.cs.siu.edu/v1"
    },
    "models": {
      "replace_with_model_id_from_v1_models": {
        "name": "SIU Qwen 3.8 27B",
        "limit": {
          "context": 32768,
          "output": 4096
        }
      }
    }
  }
}

```

The `limit.context` value tells OpenCode that the model has a 32,768-token total context window. The `limit.output` value reserves up to 4,096 tokens for a response, so OpenCode compacts conversation history before all 32,768 tokens are used. These settings do not increase the API server's capacity; they must not exceed the limits configured on the SIU gateway.

8.2 Store the credential with OpenCode

Launch OpenCode and run:

```
/connect
```

Then:

- Choose **Other**.
- Enter the provider ID `siu` exactly as written in the configuration.
- Paste the assigned API key.
- Run `/models`.
- Select `siu/YOUR\_MODEL\_ID`.

Using `/connect` keeps the real key out of the project configuration.

9. Aider

AIDER

Use Aider from the terminal

Install Aider, set OpenAI-compatible variables, and open a Git project.

1 Install
python -m pip install aider-install
aider-install

2 Configure
OPENAI_API_BASE
OPENAI_API_KEY

3 Launch
aider --model
openai/MODEL_ID

4 Review
Inspect diff
Commit changes

```
Terminal — project directory
cd my-project
git status

export OPENAI_API_BASE=https://api.cs.siu.edu/v1
export OPENAI_API_KEY=$SIU_AI_KEY

aider --model openai/$SIU_AI_MODEL

> Add input validation to app.py and write tests.

# Review the proposed diff before accepting.
```

✓ **Use Git as your safety net**
Commit clean work before starting. Aider is designed around diffs and commits, which makes changes easier to inspect and undo.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Aider installation and launch workflow.

Aider is especially useful for students working from terminals, SSH sessions, Linux servers, or Git-oriented projects.

9.1 Install Aider

```
python -m pip install aider-install
aider-install
```

9.2 Set the endpoint and key

Linux, macOS, or WSL - Aider

```
export OPENAI_API_BASE="https://api.cs.siu.edu/v1"
export OPENAI_API_KEY="$SIU_AI_KEY"
```

Windows PowerShell - Aider

```
$env:OPENAI_API_BASE = "https://api.cs.siu.edu/v1"
$env:OPENAI_API_KEY = $env:SIU_AI_KEY
```

9.3 Launch Aider in a Git project

```
cd my-project
git status
aider --model "openai/$SIU_AI_MODEL"
```



Aider may warn that it does not recognize metadata for a custom model name. Copy the included `.aider.model.metadata.json` into the project root and replace its model placeholder with the exact fully qualified name used at launch, such as `openai/Qwen3.8-27B`. The template declares a 32,768-token context window and 4,096-token output limit. Aider uses this metadata for reporting; the SIU gateway still enforces the actual limits.

Included templates:

- `examples/aider/.env.example`
- `examples/aider/.aider.conf.yml`
- `examples/aider/.aider.model.metadata.json`

The key intentionally does not appear in `.aider.conf.yml`.

10. Postman



Postman collection configuration and successful response.

Postman is useful for learning REST APIs and testing requests without writing a program.

10.1 Import the supplied files

In Postman:

- Select **Import**.

- Import `examples/postman/SIU_AI_API.postman_collection.json``.
- Import `examples/postman/SIU_AI_Local.postman_environment.json``.
- Select the **SIU AI - Local Values** environment.

10.2 Set local values

Set these environment variables:

Variable	Current value
<code>`base_url`</code>	<code>`<a href="https://api.cs.siu.edu/v1">https://api.cs.siu.edu/v1</a>`</code>
<code>`api_key`</code>	Your assigned key; mark it secret where supported
<code>`model_id`</code>	Exact model ID from the first request

Do not export or synchronize an environment after placing a real key in it unless the destination is approved. Prefer Postman's local secret or vault features when available.

10.3 Run the collection

- Send **1 - List available models**.
- Copy the Qwen model's ``id`` into ``model_id``.
- Send **2 - Chat completion**.
- Confirm ``200 OK`` and the two included tests pass.

You can also import any cURL example from this guide directly through Postman's import dialog.

11. LangChain and LangGraph

FRAMEWORKS

Use LangChain and LangGraph

Connect ChatOpenAI to the gateway, then reuse it inside a graph node.

```
langchain/quickstart.py

import os
from langchain_openai import ChatOpenAI


llm = ChatOpenAI(
    base_url="https://api.cs.siu.edu/v1",
    api_key=os.environ["SIU_AI_KEY"],
    model=os.environ["SIU_AI_MODEL"],
    temperature=0.2,
)

reply = llm.invoke(
    "What problem does a database index solve?"
)
print(reply.content)

# The same llm object can be called from a LangGraph node.
```

Minimal graph

One state, one model node, one answer.



Examples included

Both LangChain and LangGraph scripts are ready to run.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

LangChain configuration and minimal LangGraph flow.

LangChain's `ChatOpenAI` integration accepts a custom OpenAI-compatible `base\_url`.

11.1 Install the packages

LangChain 1.x and LangGraph 1.x require Python 3.10 or newer. Python 3.12 is recommended for a new environment. Confirm the interpreter version before creating the virtual environment.

Linux, macOS, or WSL - LangChain

```
cd examples/langchain
python3.12 --version
python3.12 -m venv .venv
source .venv/bin/activate
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

Windows PowerShell - LangChain

```
cd examples\langchain
py -3.12 --version
py -3.12 -m venv .venv
.\.venv\Scripts\Activate.ps1
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

A virtual environment remains tied to the Python interpreter that created it. If an existing `.venv` reports Python 3.9 or older, create a new environment with Python 3.10 or newer instead of reusing it.

11.2 Run the LangChain example

```
python quickstart.py
```

Essential configuration:

```
import os
from langchain_openai import ChatOpenAI

model = ChatOpenAI(
    base_url="https://api.cs.siu.edu/v1",
    api_key=os.environ["SIU_AI_KEY"],
    model=os.environ["SIU_AI_MODEL"],
    temperature=0.2,
)

response = model.invoke("What problem does a database index solve?")
print(response.content)
```

11.3 Run the minimal LangGraph

```
python langgraph_example.py
```

The example creates a simple graph:

```
START → answer node → END
```

The same `ChatOpenAI` object is called from the `answer` node. Build this minimal version successfully before adding tools, persistence, retrieval, or multiple agents.

Included files:

- `examples/langchain/quickstart.py`
- `examples/langchain/langgraph_example.py`

12. Claude Code through LiteLLM (advanced)

ADVANCED / VERIFY FIRST

Route Claude Code through LiteLLM

LiteLLM can translate Anthropic Messages to a non-Anthropic model; validate before broad support.



Compatibility caveat

LiteLLM documents this translation path, but Anthropic does not support non-Claude models behind third-party gateways. Tool behavior may differ from native Claude.

1

Use gateway root

Set ANTHROPIC_BASE_URL without /v1.

2

Use bearer token

Set ANTHROPIC_AUTH_TOKEN to your SIU key.

3

Select model

Launch claude --model with the exact model ID.

```
Terminal
export ANTHROPIC_BASE_URL="https://api.cs.siu.edu"
export ANTHROPIC_AUTH_TOKEN="$SIU_AI_KEY"
export ANTHROPIC_MODEL="$SIU_AI_MODEL"
export CLAUDE_CODE_ENABLE_GATEWAY_MODEL_DISCOVERY=1

claude --model "$SIU_AI_MODEL"
```

URL rule

Claude Code appends /v1/messages.

Correct

`https://api.cs.siu.edu`

Do not use

`https://api.cs.siu.edu/v1`

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Claude Code environment settings and gateway URL rule.

Compatibility warning: This advanced integration translates Claude Code's Anthropic Messages traffic to Qwen through LiteLLM. Administrators must validate `/v1/messages`, streaming, and tool calls before publishing it. Reasoning streams require LiteLLM 1.96.0 or later.

12.1 Install Claude Code

Linux or WSL:

```
curl -fsSL https://claude.ai/install.sh | bash
claude --version
```

Windows PowerShell:

```
irm https://claude.ai/install.ps1 | iex
claude --version
```

On native Windows, install Git for Windows so Claude Code's Bash tool is available. Reopen the terminal if `claude` is not recognized after installation.

12.2 Configure and launch

First set `SIU_AI_KEY` and `SIU_AI_MODEL` as described in [Section 2](#). Then run the launcher from the guide's top-level directory.

Linux or WSL:

```
examples/claude-code/start-claude-code.sh
```

Windows PowerShell:

```
& .\examples\claude-code\start-claude-code.ps1
```

The launchers validate the SIU variables, map every Claude model role to Qwen, and set a 32,768-token context, 4,096-token output limit, and `medium` reasoning effort. They keep the key in the process environment.

Do not run `/login`; restart through the launcher if Claude reports missing authentication. Validate the integration first in a disposable repository with a read-only file summary, then review the first edit manually.

12.3 Troubleshooting

Error	Action
Tried to access `claude-opus-5`	Exit and relaunch with the included script.
Unsupported `reasoning_effort high`	Relaunch; the scripts set the supported value `medium`.
`Content block is not a thinking block`	Ask the administrator to upgrade LiteLLM to 1.96.0 or later.
`/login` prompt	Confirm the SIU key is set in the same terminal, then relaunch.

13. Codex CLI (advanced)

ADVANCED / VERIFY FIRST

Configure Codex CLI

Current Codex custom providers require the OpenAI Responses protocol.



Test the endpoint first

Only publish this guide after <https://api.cs.siu.edu/v1/responses> works with the deployed model, streaming, and tool calls.

```
~/codex/config.toml

model = "MODEL_ID_FROM_V1_MODELS"
model_provider = "siu"

[model_providers.siu]
name = "SIU AI"
base_url = "https://api.cs.siu.edu/v1"
env_key = "SIU_AI_KEY"
wire_api = "responses"
```

Launch sequence

Use a custom provider ID—not openai.

- 1 Copy config.toml to ~/.codex/
- 2 export SIU_AI_KEY=...
- 3 Run codex in a Git project
- 4 Approve commands manually



Expected route
POST /v1/responses

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Codex custom provider and Responses API requirement.

Compatibility requirement: Current Codex custom model providers use the OpenAI Responses wire protocol. Do not advertise Codex support until the deployed LiteLLM version and backend pass end-to-end tests on `POST /v1/responses`, including streaming and tool calls.

13.1 Install Codex

Linux, macOS, or WSL:

```
curl -fsSL https://chatgpt.com/codex/install.sh | sh
codex --version
```

Windows PowerShell:

```
powershell -ExecutionPolicy Bypass -c "irm https://chatgpt.com/codex/install.ps1
| iex"
codex --version
```

Alternatively, install it with npm on any platform that has Node.js:

```
npm install -g @openai/codex
codex --version
```

Reopen the terminal if `codex` is not recognized after installation. See the [official Codex CLI documentation](#) for current installation options.

13.2 Test the route before configuring Codex

After the administrator confirms the feature is enabled, a minimal request resembles:

```
curl --fail-with-body \  
  "https://api.cs.siu.edu/v1/responses" \  
  -H "Authorization: Bearer $SIU_AI_KEY" \  
  -H "Content-Type: application/json" \  
  -d '{"model\":"$SIU_AI_MODEL\","input\":"Reply with the word ready.\"}'
```

A working Chat Completions route does **not** prove that the Responses route works.

13.3 Configure a custom provider

Copy `examples/codex/config.toml` to `~/.codex/config.toml` and replace the model placeholder:

```
model = "replace_with_model_id_from_v1_models"  
model_provider = "siu"  
model_context_window = 32768  
model_auto_compact_token_limit = 28000  
  
[model_providers.siu]  
name = "SIU AI"  
base_url = "https://api.cs.siu.edu/v1"  
env_key = "SIU_AI_KEY"  
wire_api = "responses"
```

The context settings tell Codex that the custom SIU model has a 32,768-token window and trigger compaction at 28,000 tokens, leaving room for a response of up to 4,096 tokens. Keep these settings at the top level. Use the custom provider ID `siu`; do not overwrite a reserved built-in provider name.

13.4 Launch

```
export SIU_AI_KEY="your_assigned_key"  
codex
```

Start in a small Git repository, keep command approval enabled, and verify that requests reach the SIU gateway.

14. Kilo Code

KILO CODE

Configure Kilo Code

Create a custom OpenAI Compatible provider and let Kilo discover the model.

Settings → Providers → Add Custom

Choose OpenAI Compatible for Chat Completions-style gateways.

Provider name

SIU AI

Provider API

OpenAI Compatible

Base URL

<https://api.cs.siu.edu/v1>

API Key

.....

Model

Select the discovered Qwen 3.8 27B model

Save Provider

What happens next

Kilo queries `/v1/models`.

GET

/v1/models

Choose Qwen model

1

Safer default

Keep auto-approve disabled.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Kilo Code provider and model discovery settings.

14.1 Add a custom provider

In Kilo Code:

- Open **Settings**.
- Open **Providers**.
- Select **Add Custom**.
- Enter a provider name such as `SIU AI`.
- Select **OpenAI Compatible** as the provider API.
- Set the Base URL to `<https://api.cs.siu.edu/v1>`.
- Paste the assigned API key.
- Select the discovered Qwen 3.8 27B model.
- In the model configuration, set **Context Window** to `32768` and **Max Output Tokens** to `4096`.
- Save the provider.

Kilo can query the provider's `/v1/models` endpoint after a valid Base URL and key are entered. If the token fields are not available in the settings screen, set the custom model's limits in `kilo.jsonc`:

```
"limit": {
  "context": 32768,
  "output": 4096
}
```

These limits enable accurate context tracking and compaction for the custom SIU model. If discovery fails, verify the same key with the cURL model-list request before changing Kilo settings.

Keep automatic command approval disabled during initial use.

15. Roo Code

The screenshot shows the 'Configure Roo Code' interface. At the top, a maroon banner contains the text 'ROO CODE' and 'Configure Roo Code' with a subtitle 'Use Roo's dedicated LiteLLM provider for model discovery.' Below this, the main configuration area is divided into two columns. The left column, titled 'Roo Code → Settings → API Provider', contains four input fields: 'API Provider' (set to 'LiteLLM'), 'LiteLLM Base URL' (set to 'https://api.cs.siu.edu'), 'API Key' (masked with dots), and 'Model' (set to 'Select discovered Qwen model'). A blue information box below these fields explains why '/v1' is not in the base URL. A 'Save' button is at the bottom right of this section. The right column, titled 'First task', lists five steps: 1. Ask: Summarize this file. 2. Confirm the answer streams. 3. Request one small edit. 4. Review the diff. 5. Only then test commands. Below the list is a 'Safety' section with 'Auto-approve OFF'.

ROO CODE

Configure Roo Code

Use Roo's dedicated LiteLLM provider for model discovery.

Roo Code → Settings → API Provider
The dedicated LiteLLM option is recommended.

API Provider: LiteLLM

LiteLLM Base URL: https://api.cs.siu.edu

API Key:

Model: Select discovered Qwen model

Why no /v1 in the Base URL?
Roo's LiteLLM provider appends /v1/model/info when it discovers models. Use the gateway root.

Save

First task
Validate chat before agent actions.

- 1 Ask: Summarize this file.
- 2 Confirm the answer streams.
- 3 Request one small edit.
- 4 Review the diff.
- 5 Only then test commands.

Safety
Auto-approve OFF

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Roo Code LiteLLM configuration and safe first task.

Roo Code has a dedicated **LiteLLM** provider. It can retrieve model information from the gateway and is the preferred first configuration.

15.1 Dedicated LiteLLM provider

Configure:

- **API Provider:** LiteLLM
- **LiteLLM Base URL:** <https://api.cs.siu.edu>
- **API Key:** Your assigned SIU AI key
- **Model:** Select the discovered Qwen model

Use the gateway root without `/v1` in this dedicated provider because Roo appends its LiteLLM model-information route. Verify that the discovered model reports a **Context Window** of `32768` and **Max Output Tokens** of `4096`. If either value is wrong, ask the administrator to correct `/v1/model/info` or use the fallback below.

15.2 Fallback when LiteLLM discovery is unavailable

Some gateways restrict `/v1/model/info`. If the dedicated provider cannot list models:

- Select **OpenAI Compatible** instead.
- Set Base URL to <https://api.cs.siu.edu/v1>.
- Paste the assigned API key.
- Enter the exact model ID manually.
- Open **Model Configuration** and set **Context Window** to `32768` and **Max Output Tokens** to `4096`.

15.3 Safe first task

Ask Roo to summarize a file, confirm streaming works, request one small edit, review the diff, and only then test command execution. Keep auto-approve off.

A concise settings reference for Cline, Kilo, and Roo is included at `examples/UI_SETTINGS_REFERENCE.md`.

16. Troubleshooting

TROUBLESHOOTING

Troubleshoot by status code

Start with cURL, then fix the smallest failing layer.

401 / 403

Authentication or access
Re-copy the key; confirm the Authorization: Bearer header and model access.

404

Wrong route or base URL
OpenAI clients use <https://api.cs.siu.edu/v1>. Claude/Roo LiteLLM use the root where noted.

400

Bad model or parameter
List /v1/models; copy the exact ID; remove optional parameters and retry.

429

Rate or concurrency limit
Wait for active work to finish, reduce parallel requests, and retry with backoff.

5xx / timeout

Gateway or backend issue
Retry once; record time, request ID, client, model, and a redacted error.

✓

Best support report
Include client/version, exact route (not the key), model ID, timestamp with time zone, status code, and the smallest reproducible request.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

Troubleshooting map by HTTP status code.

Always reproduce the problem with the cURL model-list or chat request first. This determines whether the failure belongs to the account/gateway or only to the client configuration.

401 or 403 — authentication or access

- Re-copy the key without spaces or surrounding quotation marks in a graphical client.
- Confirm the header is ``Authorization: Bearer YOUR_API_KEY``.
- Confirm the key has access to the selected model.
- Make sure the client is not sending an unrelated OpenAI, Anthropic, or saved subscription credential.
- Revoke and replace the key if it may have been exposed.

404 — wrong route or base URL

- OpenAI-compatible clients normally use ``https://api.cs.siu.edu/v1``.
- Claude Code uses ``https://api.cs.siu.edu`` as ``ANTHROPIC_BASE_URL``.
- Roo's dedicated LiteLLM provider uses the root URL; Roo's OpenAI-compatible fallback uses ``/v1``.
- Do not paste the complete ``/chat/completions`` path into a field that expects only the base URL unless that client's documentation explicitly asks for it.

400 — bad model or unsupported parameter

- Run ``GET /v1/models`` again and copy the exact ``id``.
- Remove optional parameters such as response formats, reasoning settings, image inputs, or vendor-specific options.
- Reduce the request to one user message and ``temperature``.
- Confirm the JSON is valid.

``fetch`` is not defined when running JavaScript

The JavaScript example and its pinned OpenAI SDK require Node.js 22 or newer. Check with ``node --version``. If the version is older than 22, upgrade Node.js, open a new terminal, then run ``npm install`` and ``npm start`` again. Activating or deactivating a Python virtual environment does not change the Node.js runtime.

No matching distribution for `langchain-openai`

Run `python --version` inside the active virtual environment. LangChain 1.x requires Python 3.10 or newer. Installing a newer Python does not update an existing environment; deactivate the old environment and create a new one with the newer interpreter.

```
deactivate
python3.12 -m venv .venv312
source .venv312/bin/activate
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

429 — rate, budget, or concurrency limit

- Wait for an active generation to finish.
- Avoid sending multiple simultaneous requests from the same key.
- Reduce automated retries and parallel agents.
- Retry with exponential backoff rather than an immediate loop.
- Contact the service administrator if a course workflow consistently needs a different limit.

5xx or timeout — gateway or model backend issue

- Retry once after a brief delay.
- Record the time with time zone, model ID, client, route, and redacted error.
- Do not repeatedly retry a large request.
- Test a tiny prompt with cURL.
- If cURL also fails, report the gateway issue rather than reinstalling the client.

Chat works, but an agent cannot use tools

This usually indicates a tool-calling compatibility problem rather than an authentication problem. Record:

- Client and exact version
- Model ID
- Whether the tool call appeared as ordinary text, malformed JSON, or a native tool event
- Whether streaming was enabled

- The smallest prompt that reproduces the problem

Use chat-only or edit-only mode until the integration is validated.

Support-report template

```
Client and version:
Operating system:
Date/time and time zone:
Base URL or route used:
Model ID:
HTTP status code:
Did the cURL quick test work?:
Redacted error message:
Smallest reproducible prompt/request:
```


Never include the API key in a support report.

17. Security and responsible use

SECURITY

Protect your API key and your project

Treat the key like a password; treat agent commands like code review.



YOUR KEY

Do
Secure habits

- ✓ Use environment variables or client secret stores.
- ✓ Add .env to .gitignore before creating it.
- ✓ Use your own assigned key—never share one.
- ✓ Review diffs and terminal commands.
- ✓ Revoke a key immediately if exposed.

Do not
Common mistakes

- ✗ Put the key in source code or screenshots.
- ✗ Commit .env, settings, or exported secrets.
- ✗ Paste protected university data without approval.
- ✗ Enable unrestricted auto-approve.
- ✗ Assume generated code is correct or secure.

SIU AI API • api.cs.siu.edu

Illustrated configuration map — product menus may move as software updates.

API-key security and coding-agent safety guidance.

API keys

- Use an environment variable, client credential store, or approved secret manager.
- Add `.env` to `.gitignore` before creating the file.
- Do not share a key with classmates, teammates, or instructors.

- Do not reuse one person's key for a class, lab, or public service.
- Revoke a key immediately after accidental exposure.

A useful `.gitignore` entry is:

```
.env  
.env.*  
!.env.example
```

University and research data

Do not submit protected, regulated, confidential, export-controlled, student-record, health, personnel, credential, or other restricted information unless the data owner and applicable SIU policy explicitly permit that use. The fact that a service is university-hosted does not automatically authorize every category of data.

Coding agents

Coding agents can modify files, run commands, install packages, access local configuration, and follow malicious instructions embedded in repository content. Use them as supervised assistants:

- Commit or back up clean work first.
- Read proposed changes.
- Review terminal commands before approval.
- Keep auto-approve off for consequential actions.
- Do not expose `.env`, SSH keys, browser profiles, cloud credentials, or production secrets.
- Run untrusted projects in an isolated environment.
- Test generated code and review it for security, licensing, correctness, and course requirements.

Academic use

Follow the instructor's rules for each assignment. Preserve required citations and disclose AI assistance when the course policy requires it. An available API does not override academic-integrity expectations.

18. Included files

- `examples/common/.env.example`: Common endpoint, key, and model environment template.

- ``examples/curl/list-models.sh``: Lists the models available to a key.
- ``examples/curl/chat.sh``: Sends a non-streaming chat request.
- ``examples/python/quickstart.py``: Python SDK request with error handling.
- ``examples/python/streaming.py``: Python streaming example.
- ``examples/javascript/quickstart.mjs``: Node.js SDK request.
- ``examples/javascript/streaming.mjs``: Node.js streaming example.
- ``examples/continue/config.yaml``: Continue local configuration.
- ``examples/continue/.env.example``: Continue secret-file template.
- ``examples/opencode/opencode.json``: OpenCode custom provider.
- ``examples/aider/.aider.conf.yml``: Aider endpoint and model configuration.
- ``examples/aider/.aider.model.metadata.json``: Aider context and output-limit metadata.
- ``examples/aider/.env.example``: Aider secret template.
- ``examples/postman/SIU_AI_API.postman_collection.json``: Postman collection.
- ``examples/postman/SIU_AI_Local.postman_environment.json``: Postman environment template.
- ``examples/langchain/quickstart.py``: LangChain example.
- ``examples/langchain/langgraph_example.py``: Minimal LangGraph example.
- ``examples/claude-code/start-claude-code.sh``: Advanced Claude Code launcher.
- ``examples/claude-code/start-claude-code.ps1``: Windows PowerShell Claude Code launcher.
- ``examples/claude-code/settings.json.example``: Claude Code gateway settings without a key.
- ``examples/codex/config.toml``: Advanced Codex custom provider.
- ``examples/UI_SETTINGS_REFERENCE.md``: One-page settings reference for UI clients.
- ``ADMIN_PUBLISHING_CHECKLIST.md``: Pre-publication validation for administrators.
- ``MANIFEST.json``: Machine-readable package inventory and hashes.

- ``CHECKSUMS.txt``: SHA-256 checksums for package files.

Official documentation references

These links were used to validate the configuration patterns in this guide. Third-party products change; consult their current documentation if a menu or option has moved.

- [LiteLLM — OpenAI-compatible proxy examples](#)
- [Continue — OpenAI model-provider configuration](#)
- [Cline — OpenAI Compatible provider](#)
- [OpenCode — providers and custom providers](#)
- [Aider — OpenAI-compatible APIs](#)
- [Postman — import data](#)
- [LangChain — OpenAI-compatible endpoints](#)
- [LiteLLM — Claude Code with non-Anthropic models](#)
- [Anthropic — Claude Code LLM gateways](#)
- [OpenAI — Codex configuration basics](#)
- [Kilo Code — OpenAI-compatible providers](#)
- [Roo Code — LiteLLM provider](#)